


```
#Import XGBoost, Pandas, and sklearn for the function that we will use to calculate the accur
#The accuracy is required to understand how our model is performing.
```

```
from sklearn.ensemble import AdaBoostClassifier
from sklearn import datasets
from sklearn import metrics
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np
import xgboost as xgb
from sklearn.metrics import accuracy_score
```

```
#Import the wholesale customer dataset - 1 point
```

```
from google.colab import files
uploaded = files.upload()
```

```
data_set = pd.read_csv("wholesale-data.csv",header=None)
data = data_set.iloc[1:,:]
print(data)
```

 No file chosen

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|-----|----|----|-------|-------|-------|-------|-------|------|
| 1 | 2 | 3 | 12669 | 9656 | 7561 | 214 | 2674 | 1338 |
| 2 | 2 | 3 | 7057 | 9810 | 9568 | 1762 | 3293 | 1776 |
| 3 | 2 | 3 | 6353 | 8808 | 7684 | 2405 | 3516 | 7844 |
| 4 | 1 | 3 | 13265 | 1196 | 4221 | 6404 | 507 | 1788 |
| 5 | 2 | 3 | 22615 | 5410 | 7198 | 3915 | 1777 | 5185 |
| .. | .. | .. | ... | ... | ... | ... | ... | ... |
| 436 | 1 | 3 | 29703 | 12051 | 16027 | 13135 | 182 | 2204 |
| 437 | 1 | 3 | 39228 | 1431 | 764 | 4510 | 93 | 2346 |
| 438 | 2 | 3 | 14531 | 15488 | 30243 | 437 | 14841 | 1867 |
| 439 | 1 | 3 | 10290 | 1981 | 2232 | 1038 | 168 | 2125 |
| 440 | 1 | 3 | 2787 | 1608 | 2510 | 65 | 477 | 52 |

Saved successfully!

```
#Create training and test sets - 80:20 - 1 point
```

```
X = data.iloc[:,1:]
Y = data.iloc[:,0]
```

```
X = X.astype(int)
#Y = Y.astype(int)
```

```
Y = Y.map({'1':0, '2':1})
#print(Y)
```

```
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)
```

```
#print(x_train)
#print(x_test)
print(y_test.shape)
```

```
(88,)
```

#Convert the pandas dataframe into a DMatrix, an internal data structure that is used by XGBo
- 2 points

```
xgb_x_train = xgb.DMatrix(x_train, label=y_train)
xgb_x_test = xgb.DMatrix(x_test, label=y_test)
```

DMatrix is internal data structure used by XGBoost which is optimized for both memory efficiency and training speed.

#Specify the training parameters and train the model.

```
#xgb_clf = xgb.XGBClassifier()
#print(x_train.dtypes)
#xgb_clf.fit(x_train, y_train)
```

```
param = {
    'max_depth': 3, # the maximum depth of each tree
    'eta': 0.3, # the training step for each iteration
    'silent': 1, # logging mode - quiet
    'objective': 'multi:softprob',
    'num_class': 2} # error evaluation for multiclass training # the number of classes that
num_round = 20
```

```
bst = xgb.train(param, xgb_x_train, num_round)
```

#Predict the "Channel" values of the test set using the model that we just created. - 1 point

Saved successfully!



```
print(y_pred.shape)
best_preds = np.asarray([np.argmax(line) for line in y_pred])
print(best_preds.shape)
print(best_preds)
```

```
(88, 2)
```

```
(88,)
```

```
[0 0 1 0 1 0 0 1 1 0 1 0 1 1 0 0 0 1 1 1 0 1 0 0 0 0 0 0 0 0 0 0 1 1 0
 0 1 1 0 1 0 0 1 1 1 1 0 0 0 0 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 1 1 0
 0 1 0 0 1 0 1 0 1 0 1 1 1 0]
```

#Get the accuracy of the model that we have trained for the test dataset. - 1 point

```
#accuracy = accuracy_score(y_test, y_pred)
#print(accuracy)

from sklearn.metrics import precision_score

print (precision_score(y_test, best_preds, average='macro'))

0.9226579520697167
```

✓ 0s completed at 9:29 AM



Saved successfully!

