

Report on Cardio Vascular Disease Prediction

TASK 1

1. Identify columns into nominal,categorical,continuous etc

Nominal -> id, gender, smoke, alco, cardio

Ordinal -> cholesterol, gluc

Interval -> height, weight, ap_hi, ap_lo

Ratio -> age

2. Insights of the data

Data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 13 columns):
id                500 non-null int64
age               335 non-null float64
gender            329 non-null object
height            198 non-null float64
weight            336 non-null float64
ap_hi             347 non-null float64
ap_lo             332 non-null float64
cholesterol       333 non-null object
gluc              333 non-null object
smoke             326 non-null float64
alco              335 non-null float64
active            343 non-null float64
cardio            500 non-null int64
dtypes: float64(8), int64(2), object(3)
memory usage: 45.0+ KB
None
```

Data.describe()

	id	age	height	weight	ap_hi \
count	500.000000	335.000000	198.000000	336.000000	347.000000
mean	50279.916000	19490.886567	163.934343	74.347321	128.685879
std	29913.623631	2466.702487	8.258559	14.335964	18.490176
min	38.000000	14334.000000	120.000000	45.000000	12.000000
25%	23446.500000	17988.500000	159.250000	65.000000	120.000000
50%	51913.500000	19719.000000	165.000000	72.000000	120.000000
75%	78656.000000	21597.500000	168.000000	82.000000	140.000000
max	99662.000000	23479.000000	187.000000	155.000000	190.000000

	ap_lo	smoke	alco	active	cardio
count	332.000000	326.000000	335.000000	343.000000	500.000000
mean	90.060241	0.092025	0.065672	0.813411	0.502000
std	87.396945	0.289505	0.248078	0.390150	0.500497
min	60.000000	0.000000	0.000000	0.000000	0.000000
25%	80.000000	0.000000	0.000000	1.000000	0.000000
50%	80.000000	0.000000	0.000000	1.000000	1.000000
75%	90.000000	0.000000	0.000000	1.000000	1.000000
max	100.000000	1.000000	1.000000	1.000000	1.000000

3. Number of null values in each column

```
id          0
age         165
gender      171
height      302
weight      164
ap_hi       153
ap_lo       168
cholesterol 167
gluc        167
smoke       174
alco        165
active      157
cardio       0
dtype: int64
```

4. Analysis of Patients

a. Oldest Person

Age of oldest person is 64.0

b. Youngest Person

Age of youngest person is 39.0


c. Average age of Person

Average age group is 52.91


d. Median Age

Median of age is 54.00

e. Relationship between cardio and age




f. Age group whose survival rate is largest



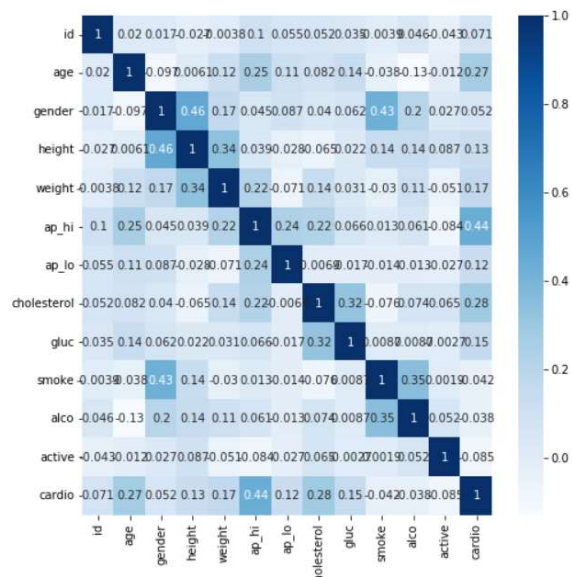
Highest survival rate is of age group 20-40 that is 81.2%


g. Similar relationship for ap_hi






Survival rate of 20-40 age group is 50.00
 Survival rate of 40-60 age group is 45.76
 Survival rate of 60-80 age group is 54.17 ----> highest

h. Visuals on data distribution







i. Missing values

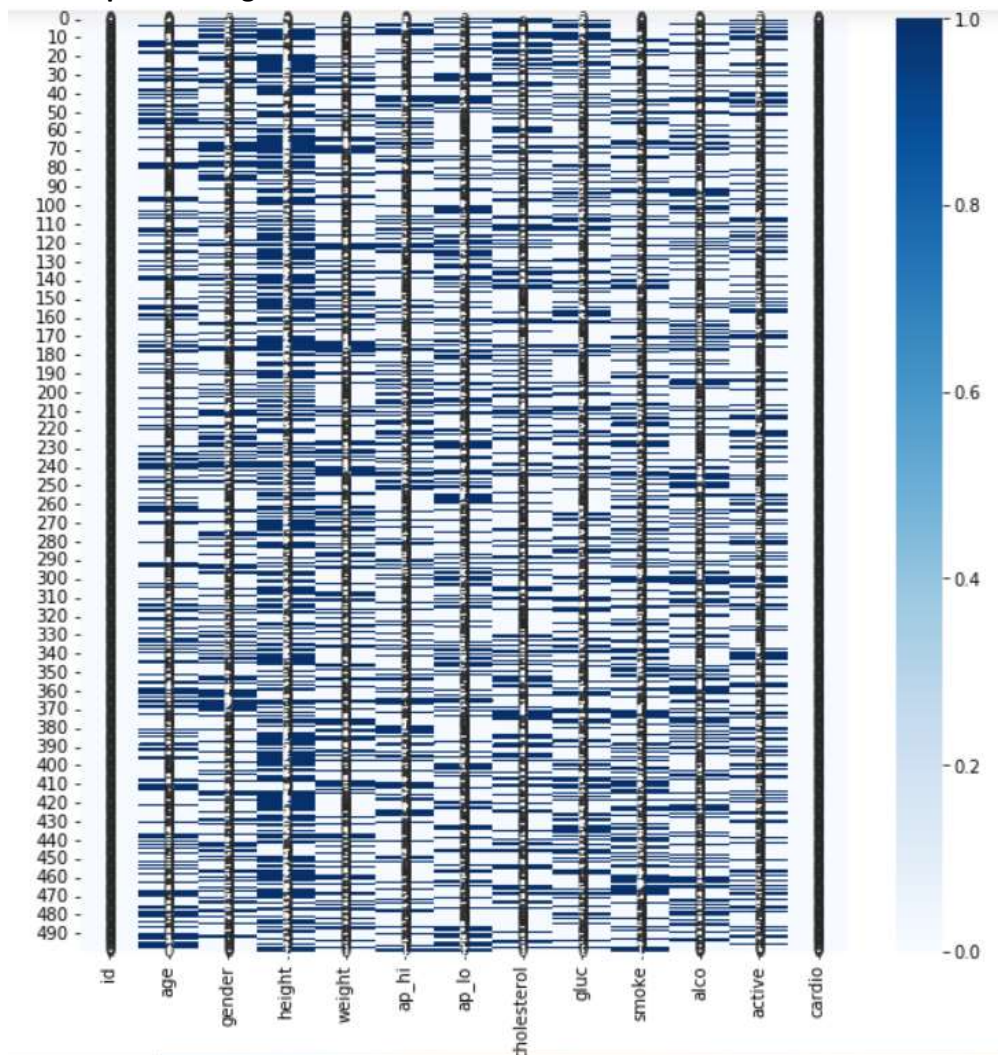
- **Count missing values**

```

id          0
age         165
gender      171
height      302
weight      164
ap_hi       153
ap_lo       168
cholesterol 167
gluc        167
smoke       174
alco        165
active      157
cardio      0
dtype: int64

```

- **Heatmap of missing values**



j. Handle missing data

- **Use Dropna**

If we drop all the rows containing any NA value only 7 rows are left, so the dimension of the returned dataframe will be (7,13)

- **Replace NA with mean**

To replace the null entries with mean of the column but for that we first need to replace null values with 0 and then find the mean for every column.

- **Additional techniques (pad,bfill)**

The columns in dataset that are not numeric wont be filled using mean values hence we need to use bfill or pad method. In bfill method the null value is replaced by next valid observation, while in pad method the null value is replaced with the previous valid observation.

k. Regression Models

- Simple Logistic Regression

Using only cardio-train.csv → Accuracy = 56.7 %

Using cardio-train.csv+cardio-validation.csv → Accuracy = 65.6 %

- Feature Selection + Logistic Regression

Using only cardio-train.csv → Accuracy = 60.8 %

Using cardio-train.csv+cardio-validation.csv → Accuracy = 61.6 %

- Random Forest

Using only cardio-train.csv → Accuracy = 68.8 %

Using cardio-train.csv+cardio-validation.csv → Accuracy = 62.4 %

- Random Forest with Grid Search CV

Using only cardio-train.csv → Accuracy = 65.6 %

Using cardio-train.csv+cardio-validation.csv → Accuracy = 69.6 %

Using cardio-train.csv+cardio-validation.csv → Accuracy = 70.4 %

l. Own implementation of Logistic Regression

```
class LR:
    def __init__(self, lr=0.001, iterations=1000):
        self.lr=lr
        self.iterations=iterations
        self.weights= None
        self.bias = None

    def fit(self,X,y):
        m, n = X.shape
        self.weights = np.zeros(n)
        self.bias=0

        for _ in range (self.iterations):
            linear_model=np.dot(X, self.weights) + self.bias
            y_pred=self.sigmoid_function(linear_model)

            dw = (1/m)*np.dot(X.T, (y_pred-y))
            db = (1/m)*np.sum(y_pred - y)

            self.weights -= self.lr*dw
            self.bias -= self.lr*db
```



```

def predict(self,X):
    linear_model = np.dot(X, self.weights) + self.bias
    y_pred = self.sigmoid_function(linear_model)
    y_pred_class = [1 if i>0.5 else 0 for i in y_pred]
    return y_pred_class

def sigmoid_function(self,x):
    return 1 / (1 + np.exp(-x))

X = train_validation[all_features]
Y = train_validation['cardio']
print(X.shape)
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.20, random_state=50)

reg_lr= LR( lr= 0.0001, iterations=1000)
reg_lr.fit(x_train, y_train)

pred = reg_lr.predict(x_test)

print("Accuracy score :", np.sum(y_test==pred) / len(y_test))

```

m. Kaggle submission

First 8 submission were done by fitting the only with cardio-train.csv. Then I trained every model with cardio-train+cardio-validation and further submissions were using this.

TASK 2

Splitting cardio-complete.csv and applying regression model on it

```
completed_data = pd.read_csv("cardio-complete.csv")
#print(completed_data.isnull().sum()) # there are no null values

completed_data.iloc[1:,2] = completed_data.iloc[1:,2].map({'Men':0, 'Women':1})
completed_data.iloc[1:,7] = completed_data.iloc[1:,7].map({'Normal':0, 'Above Normal':1, 'High':2})
completed_data.iloc[1:,8] = completed_data.iloc[1:,8].map({'Normal':0, 'Above Normal':1, 'High':2})

#print(completed_data)
X = completed_data.iloc[1:,:12]
Y = completed_data.iloc[1:,-1]

x_train, x_test, y_train, y_test = train_test_split(X,Y,test_size=0.2)

LRModel2.fit(x_train,y_train)

complete_prediction = LRModel2.predict(x_test)

print("For Task 2 ")
print('Model Accuracy : ',accuracy_score(y_test, complete_prediction))

cm = confusion_matrix(y_test, complete_prediction)
print("Confusion Matrix = ",cm)

precision = precision_score(y_test, complete_prediction,pos_label=1,labels=[0,1])
print("Precision: {:.2f}".format(precision))

recall = recall_score(y_test, complete_prediction,pos_label=1,labels=[0,1])
print("Recall: {:.2f}".format(recall))

# f1: 2 tp / (2 tp + fp + fn)
f1 = f1_score(y_test, complete_prediction,pos_label=1,labels=[0,1])
print("F1 Score is {:.2f}".format(f1))
```

TASK 3

Applying feature transformation

- **Varying polynomial degree change the accuracy**

if we choose all features ----->> for polynomial degree = 1 we are getting highest accuracy(71.5%), as we increase the degree >1 the accuracy goes on decreasing for degree=2 it was 70% for degree=3 it was 67.5%

if we choose selected features ----->> degree=4 gives maximum accuracy of 71%

- **Is the model overfitting or underfitting**

Min Cross validation Score is 0.575

Mean Cross validation Score is 0.7133262976860537

Max Cross validation Score is 0.926829268292683

The mean cross validation score is 71%, hence we can say that the model is not overfitted